

Algorithms


SF Coding Bootcamp

What is an algorithm?

- A *precise* list of instructions to solve a specific problem
 - Bad: If the list is long, break it into two
 - Good: If the list's length is longer than 100, divide the list into two equal parts
- This list of instructions must be implementable as a computer program
 - Bad: Do this if it's hot outside
- The algorithm must solve the problem *for any legal input*.
 - Example: If the algorithm finds the shortest path in a map, then it must work for any map and any source and destination.


Sorting via Insertion sort

- Let's dive right into it: The problem that we will solve is to sort a list of ordered elements in an ascending order.
- We will solve this problem using the **Insertion Sort Algorithm**.



Insertion sort: Python code

```
def insert_sorted(sorted_list, element):  
    for i, item in enumerate(sorted_list):  
        if element <= item:  
            sorted_list.insert(i, element)  
            break  
    else:  
        sorted_list.append(element)
```



```
def insertion_sort(unsorted_list):  
    sorted_list = []  
    for item in unsorted_list:  
        insert_sorted(sorted_list, item)  
    return sorted_list
```

3	1	4	7	5
---	---	---	---	---



3

1	3
---	---

1	3	4
---	---	---

1	3	4	7
---	---	---	---

1	3	4	5	7
---	---	---	---	---

Time complexity

- The most important criteria for performance of an algorithm is its *running time*. This is called the **Time Complexity** of an algorithm.
- The running time is measured the *number of primitive steps* the algorithm needs to finish its work.
 - We can consider operations like assignment, addition, comparison, etc. as primitive steps. Be careful to not assume that calls to built-in functions and methods are primitive operations.
- The running time is expressed as a function of input size.
 - The definition of input size is problem dependent. In case of sorting, the input size is the number of elements that need to be sorted.
- The number of steps can be different for same input size.
 - Usually, we consider the *longest running time* possible as our indicator of running time.

Analyze: insert_sorted()

<code>def insert_sorted(sorted_list, element):</code>	←	sorted_list has k elements.
<code> for i, item in enumerate(sorted_list):</code>	←	We loop k times. In each iteration:
<code> if element <= item:</code>	←	We do 1 comparison
<code> sorted_list.insert(i, element)</code> <code> break</code>	←	insert() is called at most once. insert() can take upto k steps by itself.
<code> else:</code>		
<code> sorted_list.append(element)</code>	←	append() is called at most once. append() can be assumed to take 1 step.

Comparison steps: Upto k

Insert: At most once, takes upto k steps

Append: At most once, takes upto 1 steps

Total steps = Comparison steps + max(Insert steps, Append steps) = $k + \max(k, 1) = \max(2 * k, 1)$

Analyze: insertion_sort()

```
def insertion_sort(unsorted_list):  
    sorted_list = []
```



unsorted_list has **n** elements.

```
    for item in unsorted_list:
```



We loop **n** times. In each iteration:

```
        insert_sorted(sorted_list, item)
```



insert_sorted() takes $2 * \text{len}(\text{sorted_list})$ steps. $\text{len}(\text{sorted_list})$ increases from 0 to **n** through the course of this loop.

```
    return sorted_list
```

1st iteration: $2 * 0$ steps (actually 1 step)

2nd iteration: $2 * 1$ steps

...

Kth iteration: $2 * k - 1$ steps

...

Nth iteration: $2 * n - 1$ steps

Total = $1 + 2 * (1 + 2 + 3 + \dots + n - 1) = 1 + n * (n - 1)$ steps

Big O notation

- In the last slide, we found that our implementation of insertion sort takes upto $1 + n * (n-1)$ steps.
- Usually we are not concerned with the exact number of steps. The most important consideration is how the running time grows as input size grows.
- We represent the worst case time in **$O(f(n))$** notation, where n is the input size, and $f()$ is a function.
 - We say an algorithm has a time complexity of $O(f(n))$, if the algorithm always takes less than **$C * f(n)$** steps, for some constant C .
- In case of insertion sort: Running time = $n^2 - n + 1$. We can write this as $O(n^2)$
- Very Very important to understand this notation. It's used by everyone talking about algorithms.

Faster Sorting: Merge Sort

- Now that we have discussed the Insertion sort algorithm and understood that its time complexity is $O(n^2)$, we can ask the question: Is there a sorting algorithm that can run faster?
- The answer is yes. In fact there are tens of such algorithms.
- We are going to look at one of them - it's called Merge sort.
- Merge sort uses an approach called **Divide and Conquer**. The idea is to divide the problem into smaller sub-parts, solve the problem for those sub-parts, and then combine the results.

Merge Sort: Merging sorted lists

- The essential element of merge sort is the ability to merge two sorted lists into one sorted list which contains all the elements of the two lists.
- The code for this is left as an exercise. We are going to call this function `merge_sorted()`.
- I leave it as another exercise for you to prove that it's time complexity is $O(m + n)$, where m and n are the sizes of each of the lists to be merged.

1	3	5
---	---	---




2	4	6
---	---	---




1	2	3	4	5	6
---	---	---	---	---	---

Merge Sort

Divide

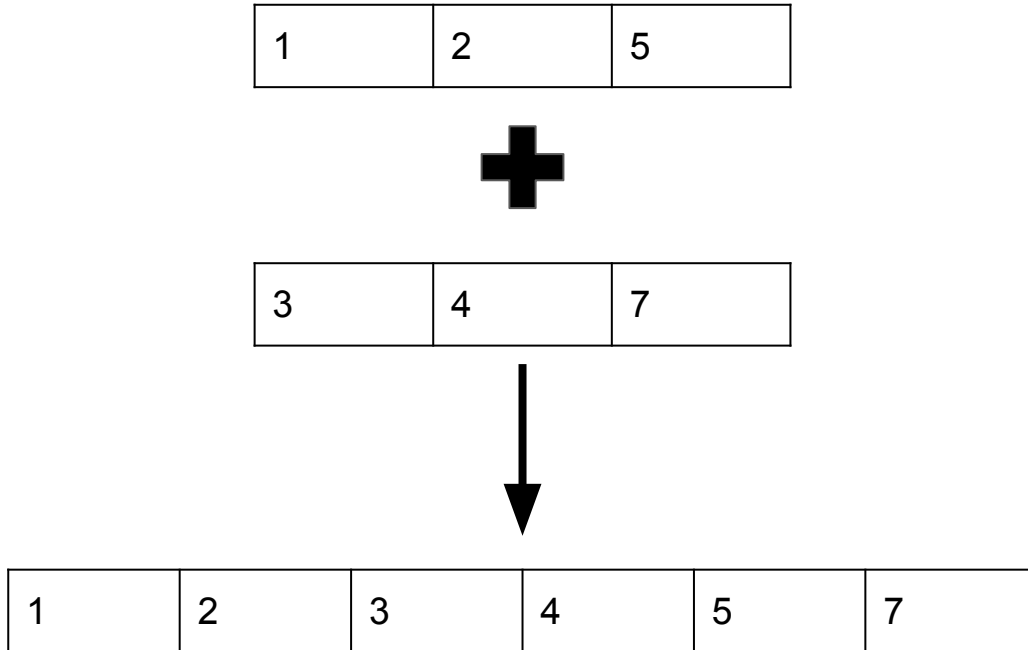


Conquer sub-parts (recursively)



Merge Sort (contd)

Merge sub-parts



Merge Sort Code

```
def merge_sort(alist):  
    if len(alist) <= 1:  
        return alist  
  
    mid = len(alist) // 2  
    left_sorted = merge_sort(alist[:mid])  
    right_sorted = merge_sort(alist[mid:])  
  
    return merge_sorted(left_sorted, right_sorted)
```

To analyze merge sort, we have to try a new approach. We will denote the Time Complexity function of Merge Sort as **$T(n)$** .

The actual computation is done on the next slide. It is important to remember that `merge_sorted()` function itself has a time complexity of $O(m + n)$.

Merge Sort Analysis

```
def merge_sort(alist):
```



alist has n elements.

```
    if len(alist) <= 1:
```



We are doing one comparison here.

```
        return alist
```

```
    mid = len(alist) // 2
```

```
    left_sorted = merge_sort(alist[:mid])
```



We are calling `merge_sort()` recursively on two lists that are half the size. Thus, these steps should take $2 * T(n/2)$ steps.

```
    right_sorted = merge_sort(alist[mid:])
```

```
    return merge_sorted(left_sorted, right_sorted)
```



This step called `merge_sorted()` function, which takes $O(n/2 + n/2)$ steps.

Total Steps: $T(n) = 1 + 2 * T(n/2) + C * n$

Merge Sort Analysis (contd)

- We have arrived at a formula to calculate the time complexity of merge sort (this is simplified): $T(n) = 2 * T(n/2) + C * n$.
- This kind of an equation is called a **recurrence relation**.
- The answer to this equation turns out to be: $T(n) = O(n * \log_2 n)$.
- As shown in a table below, this is much faster than an algorithm which has the time complexity of $O(n^2)$.

Input Size	N^2	$N * \log_2 N$
100	10,000	664
10,000	100,000,000	132,877
1,000,000	1,000,000,000,000	19,931,568

Bounds on sorting algorithms

- Considering that we were able to improve the time complexity of sorting algorithms from $O(n^2)$ to $O(n * \log_2 n)$, it's fair to ask if we can keep going.
- The answer is no: A *general purpose* sorting algorithm must take $C * n * \log_2 n$ steps in the worst case. This has been proven.
- There are other sorting algorithms (like [bucket sort](#)) which can perform better for certain kinds of inputs, and in general, the research on sorting continues to try and find better algorithms for special use cases.

Binary Search

- We are going to look at another algorithm which is very popular.
- The problem is as follows: Check if a number k is present in a sorted list of numbers. If it is present, return the position at which it is present.
- Note that the list in which we are searching is sorted. If that were not so, there is no way to optimize the solution to better than the trivial solution.
 - What is the time complexity of the trivial solution?

Binary Search solution

Find if 5 is present in the list:

1	2	3	4	5	7
---	---	---	---	---	---

Check if 5 is in first half or second half. Since $3 < 5$, 5 is in the second half. The problem becomes:

Find if 5 is present in the list:

4	5	7
---	---	---

Check if 5 is in first half or second half. Since $4 < 5$, 5 is in the second half. The problem becomes:

Find if 5 is present in the list:

5	7
---	---

Then we find 5 in the next check.

Binary Search code

```
def binary_search(k, alist, left=0, right=-1):  
    if not alist:  
        return -1  
  
    if right < left:  
        right = len(alist)  
  
    if (right - left) == 1:  
        if alist[left] == k:  
            return left  
        else:  
            return -1  
  
    mid = (left + right) // 2  
    if k < alist[mid]:  
        return binary_search(k, alist, left, mid)  
    else:  
        return binary_search(k, alist, mid, right)
```

As you can see, we check the middle element, and based on that comparison the `binary_search()` function recursively with only a subpart of the list.

In the analysis, we will again assume that the time complexity of binary search is **$T(n)$** .

Binary Search analysis

```
def binary_search(k, alist, left=0, right=-1):
```

```
    if not alist:
```

```
        return -1
```

```
    if right < left:
```

```
        right = len(alist)
```

```
    if (right - left) == 1:
```

```
        if alist[left] == k:
```

```
            return left
```

```
        else:
```

```
            return -1
```

```
    mid = (left + right) // 2
```

```
    if k < alist[mid]:
```

```
        return binary_search(k, alist, left, mid)
```

```
    else:
```

```
        return binary_search(k, alist, mid, right)
```

←

alist has **n** elements.

←

We are doing one comparison here.

←

We are doing one more comparison here.

←

We are doing another comparison here.

←

We have an assignment step and another comparison step.

←

The recursive call costs $T(n/2)$ steps.

Total Steps: $T(n) = T(n/2) + 5$

Binary Search analysis (contd)

- From the previous slide: we get the recurrence relation: $T(n) = T(n/2) + C$, where C is a constant.
- The answer to this recurrence relation is $T(n) = C * \log_2 n$. Thus, binary search has the time complexity of $O(\log_2 n)$.
- This is much faster than the trivial search algorithm which has the time complexity of $O(n)$.

Faster searching - using Python dictionary

- We just saw that if a list is sorted, it helps us perform a search inside it much faster ($O(\log_2 n)$ vs $O(n)$).
- Is there any other way we can store our data which helps us perform search even faster? As it happens, the answer is yes, and we already know about it.
- A data structure called [Hash Table](#) is designed to perform a search in $O(1)$ (i.e., constant time).
- But this is not the worst case time complexity - it is the average case time. Still, they perform very well in practice and are widely used.
- Python's dictionary data type uses a Hash Table in its implementation. That means, you can store all your data elements in a Python dictionary, and a membership check will only take $O(1)$ in the average case.